

Target-Local Soft-Information Fusion for Communication-Constrained Multi-UAV OTFS-ISAC Sensing

Xinhang Tan, Zhengquan Zhang, L Hao
Southwest Jiaotong University
Chengdu, China

Email: tanxinhang6478@163.com, zhang.zhengquan@swjtu.edu.cn, lhao@home.swjtu.edu.cn

Abstract—Multi-UAV cooperative sensing gains depend not only on which observations are selected, but also on where their soft reports are fused. We propose a target-local fusion architecture for prior-assisted OTFS-ISAC sensing. Each target is assigned to the UAV nearest its predicted position; a communication- error-calibrated, adaptive coarse-to-fine (C2F) selector then screens and replays target-dependent sensing-reporting candidates using the implemented detector. The fusion rule uses predicted beliefs only and preserves the same finite-blocklength model, 640-bit packet, 2048 channel uses, and resource caps as the canonical system. In 1000 paired trials with 15 UAVs and 10 targets, V1 attains $P_D = 0.9762$ (95% CI: [0.9730, 0.9790]) and worst-target $P_D = 0.967$. It selects 12.101 sensing observations but sends only 0.751 remote reports, corresponding to 0.481 kbit and 0.801 ms. Its paired gain over exact-marginal greedy is 0.0351, while its P_D is statistically tied with sensing-SINR selection and uses 83.7% less payload. A fusion-rule ablation and six operating- boundary sweeps show that the gain arises from target-related fusion planning jointly with two-stage selection, while large belief errors, strict rate thresholds, and sparse deployments remain the principal limitations.

Index Terms—Multi-UAV networks, OTFS-ISAC, target-local fusion, cooperative sensing, finite-blocklength reporting, coarse-to-fine selection.

I. INTRODUCTION

Integrated sensing and communication (ISAC) allows a wireless network to reuse spectrum, hardware, and signal-processing resources for transmission and environmental sensing [1], [2]. UAV-enabled ISAC further exploits mobility and favorable line-of-sight geometry to extend sensing coverage [5], [6]. When multiple UAVs observe the same target from different bistatic geometries, their soft decisions can provide complementary evidence [3].

This cooperation is not free. Each observation is produced by a sensing chain $i \rightarrow q \rightarrow j$ and must then be delivered from receiver j to a fusion UAV. A strong echo may contribute little if its reporting link is unreliable, its delay-Doppler (DD) energy is poorly captured, or its packet cost exceeds its marginal detection value [9]–[11]. Existing node-selection and backhaul-aware studies address parts of this coupling, while

distributed-detection models often abstract away the bistatic geometry and OTFS DD structure [4], [7], [8].

A further architectural choice has received less attention: *where should reports about each target meet?* A target-agnostic rule that maximizes a UAV’s aggregate incoming rate tends to assign every target to the same central receiver. That concentration changes all reporting rates, finite-blocklength success probabilities, and even the feasible candidate sets seen by the inner selector. Thus, improving only the link-ranking formula cannot correct an unfavorable fusion destination.

We study prior-assisted confirmation or re-detection and introduce target-local fusion planning. For every target, the fusion UAV is chosen using the predicted target position rather than the unknown truth. The resulting target-specific reporting graph is passed to an adaptive two-stage C2F selector whose final greedy replay is aligned with the implemented detector. This decomposition removes circular dependence between fusion planning and link selection. The distance rule is a prediction-only heuristic, not a joint optimum. The inner detector-predicted utility directly prices remote reports and independently determines how many observations to select.

Our main contributions are as follows.

- We formulate decomposed target-specific fusion planning and triplet-level soft-report selection for multi-UAV bistatic OTFS-ISAC. The fusion destination depends only on the predicted target belief, preventing truth leakage into the scheduler.
- We combine this outer planning layer with local-evidence retention, communication-error calibration, an adaptive coarse frontier, waveform-derived DD refinement, and detector-aligned greedy replay. The sensing and reporting graphs are distinct, and the detector-PD selector no longer inherits a budget from a reference algorithm.
- We provide a reproducible V1 evaluation: a 1000-trial paired main comparison, a controlled fusion-rule ablation, reporting-structure diagnostics, and sweeps over belief error, rate threshold, deployment area, communication price, blocklength, and fleet size.
- We identify the operating boundary rather than reporting a single favorable point. The results distinguish measured

serial latency from a prospective parallel-scheduling lower bound and distinguish confirmatory MC=1000 evidence from exploratory MC=100 sweeps.

The remainder of the paper presents the system model in Section II, the target-local C2F method in Section III, the V1 evidence in Section IV, and conclusions in Section V.

II. SYSTEM MODEL

We consider a sensing-centric multi-UAV OTFS-ISAC network, as shown in Fig. 1(a). The UAV and target sets are denoted by $\mathcal{M} = \{1, \dots, M\}$ and $\mathcal{Q} = \{1, \dots, Q\}$, respectively. UAV i and target q have position-velocity pairs $(\mathbf{p}_i, \mathbf{v}_i)$ and $(\mathbf{p}_q, \mathbf{v}_q)$. A cooperative candidate is indexed by the target-dependent triplet (i, j, q) , where UAV i transmits an OTFS-ISAC probing signal, target q reflects it, and UAV j receives the bistatic echo. The inter-UAV reporting topology is represented by a directed graph $\mathcal{G} = (\mathcal{M}, \mathcal{E})$, where $(i, j) \in \mathcal{E}$ if $\|\mathbf{p}_i - \mathbf{p}_j\| \leq d_{\max}$ and $i \neq j$.

We explicitly separate the *sensing* topology from the *reporting* topology. The sensing chain is $i \rightarrow q \rightarrow j$: UAV j receives the echo and therefore *produces* the soft statistic. The statistic is then reported over the link $j \rightarrow f_q$ to a target-specific fusion UAV $f_q \in \mathcal{M}$, which is the single node where all soft statistics about target q meet. Accordingly, a sensing-reporting candidate is the quadruple $(i, j, q; f_q)$, and the communication quantities that matter—reliability χ_{jf_q} , rate R_{jf_q} and delay T_{jf_q} —are properties of the reporting leg (j, f_q) alone, with no dependence on the illuminator i . Every ordered pair $i \neq j$ is a potential bistatic sensing pair; communication range constrains only the remote reporting leg. If $j = f_q$, the fusion UAV retains its locally produced statistic without transmitting a packet. This separation ensures that the fused statistic is physically well defined: with a naive $j \rightarrow i$ reporting rule, three selected links (i_1, j_1, q) , (i_2, j_2, q) , (i_3, j_3, q) would deliver their statistics to three different UAVs while the fusion step implicitly assumed they had already met. We use a geometry-driven soft-information abstraction rather than a full waveform-level OTFS receiver model. The OTFS structure enters the scheduling problem through DD-bin validity, fractional DD energy loss, same-bin collision, and local DD-grid refinement. These DD-domain quantities determine whether a target-dependent triplet provides useful soft information after reporting and fusion.

A. Communication and Reporting Model

The transmit power of UAV i is split into sensing and communication parts, i.e.,

$$P_i^s = \rho P_i, \quad P_i^c = (1 - \rho) P_i. \quad (1)$$

For a directed inter-UAV link (a, b) , the communication-channel power gain is

$$g_{ab} = \left(\frac{\lambda}{4\pi} \right)^2 d_{ab}^{-\alpha_c} \zeta_{ab} |h_{ab}^c|^2, \quad d_{ab} = \|\mathbf{p}_a - \mathbf{p}_b\|. \quad (2)$$

Here λ , α_c , ζ_{ab} , and h_{ab}^c denote the carrier wavelength, path-loss exponent, shadowing, and small-scale fading, respectively.

The conference release uses orthogonal inter-UAV reporting. The communication SINR from UAV i to UAV j is therefore

$$\gamma_{ij}^c = \frac{P_i^c g_{ij}^c}{N_0 + \epsilon_s \sum_{k \in \mathcal{M} \setminus \{i, j\}} P_k^s g_{kj}^c + \epsilon_d P_i^s g_{ij}^c}. \quad (3)$$

The coefficients ϵ_s and ϵ_d represent neighbor-induced sensing-to-communication leakage and intra-link leakage, respectively. Other report payloads do not enter (3), because strictly serial reports cannot interfere with one another.

Shared-spectrum coupling with an orthogonal reporting phase. All UAVs continuously radiate their sensing components over the same propagation field. Soft reports are then exchanged in orthogonal slots. Thus, the communication receiver sees sensing-waveform leakage, while the sensing receiver is limited by residual direct paths from the concurrently radiated sensing waveforms:

$$I_{ij}^c = \epsilon_s \sum_{k \in \mathcal{M} \setminus \{i, j\}} P_k^s g_{kj}^c, \quad (4)$$

$$I_j^s = \kappa_{\text{self}} P_j + \kappa_{\text{dc}} \sum_{k \in \mathcal{M} \setminus \{j\}} P_k^s g_{kj}^c.$$

The phase separation removes circular payload interference without removing ISAC coupling: sensing waveforms still leak into the report receiver and their residual direct paths limit sensing. Here κ_{dc} denotes the residual after reconstructing and cancelling the known illuminator waveforms.

With communication bandwidth $B_c = N\Delta f$, the reporting rate is

$$R_{ij} = B_c \log_2(1 + \gamma_{ij}^c), \quad (5)$$

and a remote reporting link is feasible when $R_{ij} \geq R_{\min}$ and $\chi_{ij} \geq \chi_{\min}$. For local evidence we define $T_{jf_q} = B_{jf_q} = 0$ and $\chi_{jf_q} = 1$ when $j = f_q$.

Finite-blocklength reliability. A soft sensing report is a short control packet. The canonical simulation uses the normal approximation

$$\epsilon_{ij} \approx Q \left(\frac{C(\gamma_{ij}^c) - \frac{k}{n}}{\sqrt{V(\gamma_{ij}^c)/n}} \right), \quad \chi_{ij} = 1 - \epsilon_{ij}, \quad (6)$$

where $C(\gamma) = \log_2(1 + \gamma)$, $V(\gamma) = (1 - (1 + \gamma)^{-2}) (\log_2 e)^2$, $k = B_q$ is the payload in bits, and n is the blocklength. Reliability, payload, and latency are then coupled through $\chi_{ij} = 1 - \epsilon_{ij}$ and $T_{ij} = n/B_c$. The heuristic reliability map is retained only for legacy-result reproduction.


B. DD-Domain Bistatic Sensing Model

For triplet (i, j, q) , define $d_{iq} = \|\mathbf{p}_q - \mathbf{p}_i\|$, $d_{jq} = \|\mathbf{p}_q - \mathbf{p}_j\|$, $\mathbf{u}_{iq} = (\mathbf{p}_q - \mathbf{p}_i)/d_{iq}$, and $\mathbf{u}_{jq} = (\mathbf{p}_q - \mathbf{p}_j)/d_{jq}$. The bistatic delay and Doppler shift are


$$\tau_{ijq} = \frac{d_{iq} + d_{jq}}{c}, \quad (7)$$

$$\nu_{ijq} = \frac{(\mathbf{v}_q - \mathbf{v}_i)^T \mathbf{u}_{iq}}{\lambda} + \frac{(\mathbf{v}_q - \mathbf{v}_j)^T \mathbf{u}_{jq}}{\lambda}.$$

(a) Multi-UAV OTFS-ISAC Scenario



(b) C2F DD-Aware Selection and Fusion



Coarse table: screening; fine table: final selection, fusion, and detection.

Fig. 1. Multi-UAV OTFS-ISAC cooperative sensing and target-specific reporting framework. Each target's reports meet at one fusion UAV selected from the predicted target belief.

The corresponding DD-grid indices are

$$l_{ijq} = \text{round}(\tau_{ijq} L \Delta f), \quad k_{ijq} = \text{round}(\nu_{ijq} N T). \quad (8)$$

A triplet is DD-valid if $0 \leq l_{ijq} < L$ and $-N/2 \leq k_{ijq} < N/2$. The fractional DD offsets are

$$\varepsilon_{ijq}^\tau = \tau_{ijq} L \Delta f - l_{ijq}, \quad \varepsilon_{ijq}^\nu = \nu_{ijq} N T - k_{ijq}. \quad (9)$$

Using the normalized sinc function, the coarse DD-domain gain is

$$\eta_{ijq}^c = \text{sinc}^2(\varepsilon_{ijq}^\tau) \text{sinc}^2(\varepsilon_{ijq}^\nu). \quad (10)$$

For C2F DD refinement, a local $(2W + 1) \times (2W + 1)$ DD neighborhood provides the waveform-derived recoverable energy η_{ijq}^{loc} . The refined gain is

$$\eta_{ijq}^f = \eta_{ijq}^{\text{loc}}. \quad (11)$$

Thus, C2F changes computational effort rather than introducing a hand-set sensing-gain interpolation.

If multiple DD-valid targets observed by the same UAV pair fall into the same DD bin, we use the collision penalty

$$\eta_{ijq}^{\text{col}} = \frac{1}{n_{ijq}^{\text{col}}}, \quad (12)$$

where n_{ijq} is the number of targets sharing bin (l_{ijq}, k_{ijq}) .

The bistatic target gain is

$$g_{ijq}^s = \frac{\lambda^2 \sigma_q}{(4\pi)^3 d_{iq}^2 d_{jq}^2}, \quad (13)$$

where σ_q denotes the target RCS. Combining bistatic gain, processing gain, DD-domain gain, and collision penalty gives

$$\gamma_{ijq}^s = \frac{\tilde{P}_i^s g_{ijq}^s G_p \eta_{ijq}^{\text{DD}} \eta_{ijq}^{\text{col}}}{N_0 + I_j^s}, \quad \eta_{ijq}^{\text{DD}} \in \{\eta_{ijq}^c, \eta_{ijq}^f\}. \quad (14)$$

In the sensing-only power model, $\tilde{P}_i^s = P_i^s$ and $G_p = NL$. Both receivers use the shared propagation field in (4), with receiver-specific suppression. The illuminator's direct path remains in I_j^s , which makes (14) near-far limited rather than interference free.

C. Soft-Information Fusion Model

Each feasible triplet produces a waveform-derived local log-likelihood ratio (LLR). After OTFS matched filtering, the delay-Doppler bin for target q on link $i \rightarrow q \rightarrow j$ is a single complex sample under an independent fast-fluctuation, Swerling-II-like model. The normalised bin energy $x = |z|^2 / \sigma_n^2$ is therefore exponential, $x \sim \text{Exp}(1)$ under $\mathcal{H}_0$ and $x \sim \text{Exp}(1 + \gamma_{ijq}^s)$ under $\mathcal{H}_1$. The local log-likelihood ratio, centred under $\mathcal{H}_0$, is

$$\ell_{ijq} = \frac{\gamma_{ijq}^s}{1 + \gamma_{ijq}^s} (x - L), \quad (15)$$

where L independent frames are incoherently integrated in one CPI, so their summed energy is Gamma distributed. Its

moments are

$$\begin{aligned}\delta_{ijq} &= \mathbb{E}_1[\ell] - \mathbb{E}_0[\ell] = \frac{L(\gamma_{ijq}^s)^2}{1 + \gamma_{ijq}^s}, \\ \sigma_{ij,0}^2 &= \text{Var}_0[\ell] = \frac{L(\gamma_{ijq}^s)^2}{(1 + \gamma_{ijq}^s)^2},\end{aligned}\quad (16)$$

The LLR carries no free scale parameter, and its single-link deflection $\delta_{ijq}^2/\sigma_{ij,0}^2 = L(\gamma_{ijq}^s)^2$ follows directly from this independent-look model.

Communication-error calibration acts on the soft statistic through the reporting leg $j \rightarrow f_q$:

$$\delta_{ijq}^{\text{eff}} = \chi_{jf_q} \delta_{ijq}, \quad \bar{\sigma}_{ij,0}^2 = \chi_{jf_q} \sigma_{ij,0}^2 + (1 - \chi_{jf_q}) \sigma_{\text{err}}^2, \quad (17)$$

where χ_{jf_q} is the success probability (6) of the report $j \rightarrow f_q$. The same post-channel moments generate the selection metric, the detection threshold, and the Monte Carlo samples.

For selected set S_q , the released V1 model assumes conditionally independent local statistics, giving

$$D_q(S_q) = \sum_{(i,j) \in S_q} \frac{(\delta_{ijq}^{\text{eff}})^2}{\bar{\sigma}_{ij,0}^2}. \quad (18)$$

Modeling correlated reports is left to future work.

Using the Gaussian deflection approximation, the detection probability is

$$P_{D,q} \approx Q\left(Q^{-1}(P_{\text{FA}}) - \sqrt{D_q(S_q)}\right), \quad (19)$$

where $Q(\cdot)$ denotes the Gaussian Q -function. Because the local LLR has a positively skewed centred-Gamma null distribution, the Monte Carlo detector calibrates the threshold with the Cornish–Fisher correction $z_{\text{CF}} = z + (\gamma_{1,0}/6)(z^2 - 1)$, where $\gamma_{1,0}$ is computed from the same post-report moments used by fusion.

D. Truth, Belief and the Scheduling Loop

We study prior-assisted confirmation or re-detection after a coarse acquisition stage; blind first acquisition is outside the present scope. The scheduler therefore receives an upstream predicted belief but cannot access the true target state. We carry two states side by side. The *truth* ($\mathbf{p}_{q,t}, \mathbf{v}_{q,t}$) drives echo generation, the true delay-Doppler, the true sensing gain and the detector; the *belief* $b_{q,t} = \mathcal{N}(\hat{\mathbf{x}}_{q,t}, \mathbf{P}_{q,t})$ drives link selection and the DD search window. The intended tracking loop is

$$\begin{aligned}\text{predict} &\rightarrow \text{schedule} \rightarrow \text{sense} \rightarrow \text{update}, \\ \hat{\mathbf{x}}_{t|t-1} &= \mathbf{F} \hat{\mathbf{x}}_{t-1|t-1}, \\ \mathbf{P}_{t|t-1} &= \mathbf{F} \mathbf{P}_{t-1|t-1} \mathbf{F}^T + \mathbf{Q}.\end{aligned}\quad (20)$$

The simulator instantiates this belief as an upstream Gaussian snapshot rather than implementing the posterior update. Its covariance is propagated through the bistatic DD Jacobian to form a three-sigma gate. Scheduling uses mean RCS; the fast look-to-look fluctuation is marginalized by (15). Thus, neither target-state truth nor a realized scattering coefficient enters scheduling.

III. TARGET-LOCAL ADAPTIVE C2F SELECTION

A. Prediction-Only Fusion Planning

Let $\hat{\mathbf{p}}_q$ be the predicted position supplied by the upstream tracker. V1 assigns target q to

$$f_q = \arg \min_{m \in \mathcal{M}} \|\hat{\mathbf{p}}_q - \mathbf{p}_m^{\text{UAV}}\|_2. \quad (21)$$

The assignment is recomputed from the predicted belief; the true target state and current-CPI RCS realization are unavailable to the scheduler. Unlike the previous target-agnostic rule $\arg \max_m \sum_j R_{jm}$, (21) can map different targets to different fusion UAVs. It therefore changes the reporting quantities R_{jf_q} and χ_{jf_q} before link selection begins.

For target q , the resulting feasible candidate set is

$$\begin{aligned}\mathcal{E}_q &= \{(i, j) : (l_{ijq}, k_{ijq}) \text{ is DD-valid}, \\ &\quad j = f_q \text{ or } (j, f_q) \in \mathcal{E}_c\}.\end{aligned}\quad (22)$$

For a remote edge, $R_{jf_q} \geq R_{\min}$ and $\chi_{jf_q} \geq \chi_{\min}$ are also required. If $j = f_q$, the statistic is already local: $B_e = T_e = 0$ and $\chi_e = 1$. Otherwise, the released serial finite-blocklength MAC charges 640 bits and n/B_c seconds per report.

B. Fair Detector-Aligned Marginal

Let $\hat{\mathbf{P}}_D(\mathcal{S})$ denote the detector-predicted target probabilities obtained from the same post-channel moments and threshold used in evaluation. Define $\bar{P}_{D,q} = \min(\hat{P}_{D,q}, P_D^{\text{req}})$. We use the weak-target potential

$$\begin{aligned}U(\hat{\mathbf{P}}_D) &= -Q\tau_\alpha \log \sum_{q \in \mathcal{Q}} \exp[-\bar{P}_{D,q}/\tau_\alpha] \\ &\quad - \frac{\mu_\alpha}{2P_D^{\text{req}}} \sum_{q \in \mathcal{Q}} [P_D^{\text{req}} - \hat{P}_{D,q}]_+^2.\end{aligned}\quad (23)$$

Saturation at P_D^{req} prevents the selector from spending observation or communication resources on numerically tiny gains after a target has reached the declared design point.

For candidate $e = (i, j, q)$ and current selection $\mathcal{S}$, the implemented greedy replay evaluates

$$\Delta \mathcal{L}_e = U(\hat{\mathbf{P}}_D(\mathcal{S} \cup \{e\})) - U(\hat{\mathbf{P}}_D(\mathcal{S})) - \lambda_c c_e, \quad (24)$$

where $c_e = \mathbf{1}\{j \neq f_q\}n/B_c$ is the per-report communication price in milliseconds. Because n is fixed, this is a remote-report cardinality price rather than a link-dependent delay optimizer. The replay uses the same post-channel moments and threshold correction as the Monte-Carlo detector. Starting from $\mathcal{S} = \emptyset$, it commits the largest positive marginal while enforcing $|\mathcal{S}_q| \leq L_{\max}$ and $|\mathcal{S}| \leq L_{\text{tot}}$.

C. Adaptive Two-Stage Refinement

Computing waveform-derived local DD energy for every feasible triplet is unnecessary. The coarse stage evaluates candidates using η_{ijq}^c and maintains a dynamic frontier. Promising candidates are refined with $\eta_{ijq}^f = \eta_{ijq}^{\text{loc}}$, after each commit, affected marginals are refreshed and the fine stage replays the exact detector-aligned rule in (24). This is an adaptive shortlist rather than a fixed one-shot top- K filter. In the V1

TABLE I
TARGET-LOCAL FUSION V1 SIMULATION PARAMETERS

Parameter	Value
UAVs / targets	$M = 15, Q = 10$
Area / communication range	$4000^2 \text{ m}^2, d_{\max} = 2500 \text{ m}$
UAV / target speed	$30\text{--}60 / 50\text{--}90 \text{ m/s}$
Carrier / DD grid	$5.9 \text{ GHz}, N = L = 64, \Delta f = 30 \text{ kHz}$
Power split / transmit power	$\rho = 0.8, P_i = 1 \text{ W}$
Rate threshold / false alarm	$R_{\min} = 0.2 \text{ Mbit/s}, P_{\text{FA}} = 0.05$
Report / blocklength	$640 \text{ bits}, n = 2048 \text{ channel uses}$
Predicted belief	$\sigma_p = 150 \text{ m}, \sigma_v = 15 \text{ m/s}$
Selection	$P_D^{\text{req}} = 0.95, \lambda_c = 0.005, L_{\max} = 6, L_{\text{tot}} = 60$
Fusion / reporting	nearest predicted target / serial orthogonal MAC
C2F refinement	$W = 1$, adaptive frontier, detector-aligned replay
Main / sweep trials	1000 / 100 per condition

main experiment, the procedure evaluates 117.7 local DD neighborhoods per trial, versus 509.3 under full refinement, a 76.9% reduction.

The implemented task objective is

$$F(\mathcal{S}) = U(\hat{\mathbf{P}}_D(\mathcal{S})) - \lambda_c \sum_{e \in \mathcal{S}} c_e. \quad (25)$$

The cost term can make F non-monotone, while the soft minimum couples targets and the constraints combine per-target and total caps. We therefore make no universal submodular approximation claim. V1 is an architectural extension of the existing selector: its principal new variable is f_q in (21), not a new greedy score.

IV. SIMULATION RESULTS

A. Protocol and Nominal Configuration

We evaluate prior-assisted confirmation in a $4000 \text{ m} \times 4000 \text{ m}$ region with $M = 15$ UAVs and $Q = 10$ targets. Table I lists the released V1 configuration. All main methods use the same paired geometries, detector randomness, physical model, and observation caps. The confirmatory comparison uses 1000 Monte-Carlo trials. Fusion and operating-envelope sweeps use 100 trials per condition with seed 2026 and are therefore diagnostic rather than substitutes for the main evidence.

B. Confirmatory Main Comparison

Figure 2 compares target-local adaptive C2F with two strong rules under common caps. Exact-marginal greedy uses the deflection-based fair utility without detector-aligned C2F replay; Sensing-SINR ranks nominal sensing strength and matches V1's per-target observation counts. V1 attains $P_D = 0.9762$ (95% CI: $[0.9730, 0.9790]$), versus 0.9411 for exact-marginal greedy and 0.9783 for Sensing-SINR. Worst-target values are 0.967, 0.928, and 0.970. The paired V1-minus-baseline intervals are $+0.0351$ ($[0.0310, 0.0392]$) and -0.0021

($[-0.0051, 0.0009]$), respectively. Thus, V1 improves on exact-marginal selection but is statistically tied with Sensing-SINR in detection probability.

The communication comparison is not tied. V1 selects 12.101 observations and sends 0.751 remote reports on average, corresponding to 0.481 kbit and 0.801 ms. Sensing-SINR selects the same 12.101 observations but sends 4.595 reports, 2.941 kbit, and 4.901 ms. V1 therefore reduces payload and serial delay by 83.7% without a significant P_D change. A separate MC=200 control applies the same detector-PD greedy rule after refining all candidates. Adaptive C2F uses 61.06 rather than 853.65 fine evaluations (92.85% fewer) and attains $P_D = 0.980$ versus 0.973 for full refinement. The paired difference is 0.0070 ($[0.0029, 0.0111]$), consistent with shortlist regularization under belief mismatch; full refinement is therefore a computational control, not an oracle.

C. Why Target-Local Fusion Helps

Figure 3 changes only the fusion destination rule. The target-agnostic max-in-rate rule gives $P_D = 0.919$, worst-target $P_D = 0.870$, 6.65 remote reports, and 7.093 ms. A max-min-rate alternative gives 0.969, 0.940, 6.40, and 6.827 ms. Nearest-target fusion reaches 0.982 and 0.950 with only 0.74 remote reports and 0.789 ms. It also reduces selected observations from 19.89 and 15.61 to 12.24. The mechanism is stronger than short-link selection alone: a fusion UAV can retain every statistic it produces locally, with unit reliability and no packet cost.

The reporting structure explains the difference. Max-in-rate and max-min-rate assign all ten targets to one fusion UAV on average. Nearest-target assigns them across 7.13 UAVs and lowers the mean conflict-graph requirement from 6.65/6.40 to 0.70 slots. Under nearest-target fusion, many trials need no remote slot at all. V1 still reports measured serial latency; conflict slots are retained only as a structural diagnostic.

D. Operating Envelope and Design Choices

Figure 4 summarizes six controlled scans. First, V1 remains near $P_D = 0.99$ for position-error standard deviations up to 50 m, then decreases to 0.982 at 150 m, 0.937 at 300 m, and 0.900 at 500 m. Belief capture remains 0.991 at 500 m; the degradation is therefore more consistent with prediction-induced fusion and link mismatch than with DD-gate misses alone. Second, changing $R_{\min}$ from 0.1 to 2 Mbit/s leaves $P_D = 0.982$ and all targets feasible because local observations bypass the reporting threshold. This insensitivity is specific to the present target-local architecture.

Topology is the next boundary. Increasing the square side from 2500 to 5500 m changes P_D from 0.992 to 0.932 and reports from 0.05 to 2.67. Increasing the fleet from 10 to 20 UAVs at $Q = 10$ changes P_D from 0.955 to 0.984 while reducing reports from 2.13 to 0.22. These scans vary density implicitly and must not be combined into an independent density effect.

The communication price exhibits a broad plateau. Setting $\lambda_c = 0$ requires 4.96 reports at $P_D = 0.983$, whereas




Fig. 2. Confirmatory MC=1000 comparison under a common scenario and caps. (a) Mean detection probability with 95% Wilson intervals; the dashed line marks 0.95. (b) Worst-target detection probability. (c) Paired V1-minus- baseline differences with 95% confidence intervals.




Fig. 3. Fusion-rule mechanism ablation (MC=100 per rule) with the inner adaptive C2F selector fixed. (a) Detection probability and 95% intervals. (b) Selected reports and measured serial delay. (c) Number of assigned fusion UAVs and conflict-graph slots. The slot count diagnoses parallelization potential; it is not an achieved latency.

$\lambda_c = 0.005$ uses 0.74 reports at $P_D = 0.982$. Values from 0.0025 to 0.02 have overlapping MC=100 uncertainty, so 0.005 is treated as a stable engineering point rather than a statistically unique optimum. Finally, 2048 channel uses yields $P_D = 0.982$ and 0.789 ms. The 1024-use setting gives $P_D = 0.981$ and 0.363 ms, while 3072 uses gives 0.981 and 1.216 ms. A paired high-MC comparison is required before changing the blocklength.

All figures are generated directly from the released CSV/JSON records. The resolved V1 configuration, source-data index, and plotting command accompany the manuscript, preventing exploratory and confirmatory results from being silently mixed.

V. CONCLUSION

This paper presented target-local fusion V1 for prior-assisted multi-UAV OTFS-ISAC sensing. The outer planner assigns each target to the UAV nearest its predicted position, while the existing communication-calibrated adaptive C2F selector performs detector-aligned greedy replay. The formulation therefore changes where reports meet while the

detector-PD objective independently chooses its observation set under fixed packet, blocklength, and resource caps.

In 1000 paired trials, V1 achieved $P_D = 0.9762$ and worst-target $P_D = 0.967$ from 12.101 observations, with only 0.751 remote reports, 0.481 kbit, and 0.801 ms of serial reporting. Its paired gain over exact-marginal greedy was 0.0351. Its detection rate was statistically tied with Sensing-SINR, while using 83.7% less payload. A controlled fusion-rule ablation showed that target-local planning disperses targets across fusion UAVs and exposes shorter, more reliable reporting edges to the inner selector. The evidence therefore supports a combined architectural conclusion: target-related fusion planning plus adaptive two-stage selection—rather than attributing the full improvement to a new greedy criterion.

V1 is a stable nominal version, not a universal endpoint. Large prediction errors and sparse deployments remain its principal boundaries. The present rate-threshold scan is weak because local evidence dominates remote reporting. Future work will add fusion capacity and sensing-acquisition costs, validate the analytic DD proxy at waveform level, model obstruction and correlated reports, and optimize link-dependent blocklengths instead of a fixed report price.




Fig. 4. Exploratory V1 operating envelope (MC=100 per condition). Blue circles show P_D ; orange squares show the named secondary metric. (a) Position-belief error, (b) minimum reporting rate, (c) deployment side length, (d) communication price, (e) blocklength, and (f) UAV fleet size.

REFERENCES

- [1] F. Liu *et al.*, "Integrated sensing and communications: Toward dual-functional wireless networks for 6G and beyond," *IEEE J. Sel. Areas Commun.*, vol. 40, no. 6, pp. 1728–1767, Jun. 2022, doi: 10.1109/JSAC.2022.3156632.
- [2] N. Gonzalez-Prelcic *et al.*, "The integrated sensing and communication revolution for 6G: Vision, techniques, and applications," *Proc. IEEE*, vol. 112, no. 7, pp. 676–723, Jul. 2024, doi: 10.1109/JPROC.2024.3397609.
- [3] K. Meng, C. Masouros, A. P. Petropulu, and L. Hanzo, "Cooperative ISAC networks: Performance analysis, scaling laws, and optimization," *IEEE Trans. Wireless Commun.*, vol. 24, no. 2, pp. 877–892, Feb. 2025, doi: 10.1109/TWC.2024.3491356.
- [4] K. Meng, C. Masouros, G. Chen, and F. Liu, "Network-level integrated sensing and communication: Interference management and BS coordination using stochastic geometry," *IEEE Trans. Wireless Commun.*, vol. 23, no. 12, pp. 19365–19381, Dec. 2024, doi: 10.1109/TWC.2024.3483031.
- [5] J. Mu, R. Zhang, Y. Cui, N. Gao, and X. Jing, "UAV meets integrated sensing and communication: Challenges and future directions," *IEEE Commun. Mag.*, vol. 61, no. 5, pp. 62–67, May 2023, doi: 10.1109/MCOM.008.2200510.
- [6] K. Meng *et al.*, "UAV-enabled integrated sensing and communication: Opportunities and challenges," *IEEE Wireless Commun.*, vol. 31, no. 2, pp. 97–104, Apr. 2024, doi: 10.1109/MWC.131.2200442.
- [7] A. Chawla, A. Patel, A. K. Jagannatham, and P. K. Varshney, "Distributed detection in massive MIMO wireless sensor networks under perfect and imperfect CSI," *IEEE Trans. Signal Process.*, vol. 67, no. 15, pp. 4055–4068, Aug. 2019, doi: 10.1109/TSP.2019.2924588.
- [8] A. Chawla, A. S. Sarode, A. K. Jagannatham, and L. Hanzo, "Distributed parameter detection in massive MIMO wireless sensor networks relying on imperfect CSI," *IEEE Trans. Wireless Commun.*, vol. 20, no. 1, pp. 506–519, Jan. 2021, doi: 10.1109/TWC.2020.3025877.
- [9] M. Chen, M.-M. Zhao, A. Liu, M. Li, and Q. Shi, "Joint node selection and resource allocation optimization for cooperative sensing with a shared wireless backhaul," *IEEE Trans. Signal Process.*, vol. 73, pp. 67–82, 2025, doi: 10.1109/TSP.2024.3516709.
- [10] W. Yuan *et al.*, "From OTFS to DD-ISAC: Integrating sensing and communications in the delay-Doppler domain," *IEEE Wireless Commun.*, vol. 31, no. 6, pp. 152–160, Dec. 2024, doi: 10.1109/MWC.018.2300607.
- [11] M. F. Keskin *et al.*, "Integrated sensing and communications with MIMO-OTFS: ISI/ICI exploitation and delay-Doppler multiplexing," *IEEE Trans. Wireless Commun.*, vol. 23, no. 8, pp. 10229–10246, Aug. 2024, doi: 10.1109/TWC.2024.3370501.